

Kernel PCA

Two moons dataset



Note

1) We're assume mean-centered data

$$EX=0$$

$$2) \Sigma_n = \frac{1}{n} \sum_{i=1}^n x_i x_i^T \quad (\text{biased})$$

Consider a 2D dataset $x = (x_1, x_2)$

We can lift the data to 6d:

$$\phi(x) = [1, x_1, x_2, x_1^2, x_2^2, x_1 x_2] \in \mathbb{R}^6$$

Instead of applying PCA to X

We apply PCA to $\phi(x)$

$$R_\phi = \frac{1}{n} \sum_i \phi(x_i) \phi(x_i)^T \in \mathbb{R}^{l \times l}$$

Motivation: In the high dim space we can find linear subspace (principal directions) that capture the variability of our data.

Reminder:

If u is a vector in the span of $\{X\}$

We can write u as a linear combination of the vectors x_i in X

Kernel Trick

Let v be an eigenvector of R_ϕ

We can write

$$v \in \text{span } \phi(x) \Rightarrow v = \sum_{j=1}^n \alpha_j \phi(x_j)$$

$$R_\phi v = \lambda v$$

$$\underbrace{\frac{1}{n} \sum_i \phi(x_i) \phi(x_i)^T}_{R_\phi} \left(\sum_{j=1}^n \alpha_j \phi(x_j) \right) = \lambda \sum_{i=1}^n \alpha_i \phi(x_i)$$

Multiply $\phi(x_k)^T$ on the left and rearrange sums:

$$\frac{1}{n} \sum_{j=1}^n \alpha_j \sum_{i=1}^n \phi(x_k)^T \phi(x_i) \phi(x_i)^T \phi(x_j) = \lambda \sum_{i=1}^n \alpha_i \phi(x_k)^T \phi(x_i)$$

$$K_{ij} = \phi(x_i)^T \phi(x_j) \quad (\text{kernel})$$

$$\frac{1}{n} \sum_{j=1}^n \alpha_j \sum_{i=1}^n K_{ij} K_{ki} = \lambda \sum_{i=1}^n \alpha_i K_{ki}$$

$$\Rightarrow \frac{1}{n} K K \alpha = \boxed{\frac{1}{n} K^2 \alpha = \lambda K \alpha}$$

$$K(K\alpha - \lambda n \alpha) = 0$$

$\Rightarrow \alpha$ is an EV of K

Projecting onto EV v of R_f (PCA)

$$\phi(x)^T v = \sum_{j=1}^n \alpha_j \phi(x)^T \phi(x_j) = \sum_{j=1}^n K(x, x_j) \alpha_j$$

$$= K \alpha$$

$\underbrace{\hspace{1cm}}$

$$K = k(x, x_j). \quad x_j \in X$$

Kernel PCA elevates from p to $l \gg p$
and then reduce dim to d

Instead we calculate kernel matrix $(n \times n)$
and its eigen decomposition gives the "non linear"
principal components

Kernel matrix = Gram matrix

The Gram matrix of X is given by

$$G = X^T X, \quad G_{ij} = \langle x_i, x_j \rangle = x_i^T x_j$$

Properties

1) G is PSD : $u^T G u \geq 0$ for any u
(All eigenvalues of G being non-negative)

2) G is symmetric : $G_{ij} = G_{ji}$

Kernel examples

$$K(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

$$\phi(x_i) = [1, \sqrt{2} x_1, \sqrt{2} x_2, x_1 x_2, x_1^2, x_2^2] \in \mathbb{R}^6$$

$$\begin{aligned} \phi(x_i)^T \phi(x_j) &= 1 + 2x_i(1)x_j(1) + 2x_i(2)x_j(2) + \\ &+ 2x_i(1)x_j(1)x_i(2)x_j(2) + x_i(1)^2 x_j(1)^2 \\ &+ x_i(2)^2 x_j(2)^2 = (1 + x_i^T x_j)^2 \end{aligned}$$

$$K[i, j] = (1 + x_i^T x_j)^2$$

This kernel is for ϕ that is a polynomial of degree 2.

1) Polynomials of higher degree: $d > 2$

$$K_{ij} = K[i, j] = (1 + x_i^T x_j)^d$$

2) Sigmoid kernel

$$K(x_i, x_j) = \tanh(\alpha x_i^T x_j + \beta)$$

3) Gaussian / RBF kernel

$$k(x_i, x_j) = \exp \left\{ - \|x_i - x_j\|^2 / (2\sigma^2) \right\}$$

Theorem

A function $k(x_i, x_j)$ is a valid kernel if

1) k is symmetric: $k(x_i, x_j) = k(x_j, x_i) \leftarrow$ this is easy

2) k is PSD

$$\sum_i \sum_j y_i y_j k(x_i, x_j) \geq 0 \quad \forall y \in \mathbb{R}^n$$

\leftarrow harder to check

Mercer's Theorem

A function $k(x_i, x_j)$ is PSD iff it can be expressed as an inner product:

$$k(x_i, x_j) = \langle \phi(x_i), \phi(x_j) \rangle = \phi(x_i)^T \phi(x_j)$$

Let's look at the Gaussian kernel

$$\exp \left\{ -\frac{\|x_i - x_j\|^2}{2\sigma^2} \right\} = \exp \left\{ -\frac{\|x_i\|^2}{2\sigma^2} \right\} \cdot \exp \left\{ -\frac{\|x_j\|^2}{2\sigma^2} \right\} \cdot \exp \left\{ \frac{2x_i^T x_j}{2\sigma^2} \right\}$$

$$= \exp \left\{ \frac{\|x_i\|^2}{2\sigma^2} \right\} \exp \left\{ -\frac{\|x_j\|^2}{2\sigma^2} \right\} \underbrace{\sum_{k=0}^{\infty} \frac{1}{k!} (x_i^T x_j)^k / \sigma^2}_{e^x = \sum_{h=0}^{\infty} \frac{x^h}{h!} \text{ (Taylor)}}$$

Note that the high dim feature space here is infinite-dimensional

Infinite sum of PSD is still PSD

Proof of Mercer's Theorem (high level)

1) We assume K is symmetric, therefore we can apply spectral theorem

$$K = V \Lambda V^T$$

2) If K is PSD then $\lambda_i \geq 0$ (eigenvalues)

3) We can define a mapping

$$\phi(x_j) = [\sqrt{\lambda_1} v_1(j), \sqrt{\lambda_2} v_2(j), \dots, \sqrt{\lambda_n} v_n(j)]^T$$

Then by construction we get

$$K[i,j] = \langle \phi(x_i), \phi(x_j) \rangle$$

Gaussian kernel and bandwidth selection

$$K(x_i, x_j) = \exp\{-\|x_i - x_j\|^2 / 2\sigma^2\}$$

σ is a bandwidth

1) If $\sigma \gg \|x_i - x_j\| \quad \forall x_i, x_j$

$$\frac{\|x_i - x_j\|}{\sigma} \rightarrow 0 \quad K(x_i, x_j) = 1$$

2) If $\sigma \ll \|x_i - x_j\|$

$$\frac{\|x_i - x_j\|}{\sigma} \rightarrow \infty \quad k(x_i, x_j) = 0$$

3) By correct choice of σ , the Gaussian kernel preserves Locality:

we'll have high k_{ij} for x_i, x_j that are similar
and $k_{ij} \approx 0$ for x_i, x_j that are distant

4) Typical choice $\sigma = \text{median}_{x_i, x_j \in X} \{\|x_i - x_j\|\}$

- 5) small bandwidth reduces bias - average only over points that are similar to one another but it also increases the variance in noisy data

Nystrom / Out-of-sample extension

The kernel has size $n \times n$

$O(n^2)$ to calculate
 $O(n^3)$ to calculate ^Aeigenvectors } naive estimate

Can reduce computational complexity by approximation

Assumption of Nystrom approach: kernel is low-rank

(rank $d \ll n$)

Sample m random points from the data

Calculation will rely on computing K only

between the m points and the n points

$$K = \begin{bmatrix} K_{11} & K_{12} \\ K_{21} & K_{22} \end{bmatrix} \quad K_{12} = K_{21}^T \quad \leftarrow$$



$$\begin{aligned} K_{11} & m \times m \\ K_{21} & (n-m) \times m \\ K_{22} & (n-m) \times (n-m) \end{aligned}$$

For low rank approximation

$$K = U_d \Lambda_d U_d^T \quad \leftarrow$$

$$\text{where } U = \begin{bmatrix} U_1 \\ U_2 \end{bmatrix} \in \mathbb{R}^{n \times d}$$

Λ - diagonal $d \times d$ matrix

We're going to take an EVD of K_{11} ($O(m^3)$)

$$\begin{aligned} K_{11} &= U_1 \Lambda U_1^T \\ K_{21} &= U_2 \Lambda U_1^T \end{aligned} \quad \underbrace{\begin{bmatrix} U_1 \\ U_2 \end{bmatrix}}_{d \text{ eigenvectors}} \begin{bmatrix} \Lambda \\ 0 \end{bmatrix} \begin{bmatrix} U_1 & U_2 \end{bmatrix}^T$$

$$K_{21} U_1 = U_2 \Lambda$$

$$\underline{K_{21} U_1} \underline{\Lambda^{-1}} = \underline{U_2} \quad \leftarrow$$

$$K_{22} = U_2 \Lambda U_2^T$$

Summary

- 1) PCA can't "identify" non linear structures
- 2) Kernel PCA - map data to high dim space
Apply PCA there

We avoid having to calculate K_p ($l \times l$)

by calculating α Eigenvectors of $K_{ij} = \phi(x_i)^T \phi(x_j)$

$$3) \underbrace{\phi(x_i)^T v}_{\text{don't have to calculate}} = [K_{\alpha}]_i$$

- 4) All properties of PCA (max variance, minimum reconstruction error, uncorrelated features) hold in the high-dim space.

- 5) Kernel PCA just relies on eigenvale decomp
- 6) Flexible: can use different kernels

Limitations:

- 1) Reconstruction? need pre-image solution
- 2) Interpretability